

In []:

Name:Dipali Chavan

Roll No: I4139

In [1]:

```
import numpy as np
import pandas as pd
import random
import tensorflow as tf
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Flatten, Conv2D, Dense, MaxPooling2D
from tensorflow.keras.optimizers import SGD
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.datasets import mnist
```

In [2]:

```
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

In [3]:

```
print(X_train.shape)
```

(60000, 28, 28)

In [4]:

```
X_train[0].min(), X_train[0].max()
```

Out[4]:

```
(0, 255)
```

In [5]:

```
X_train = (X_train - 0.0) / (255.0 - 0.0)
X_test = (X_test - 0.0) / (255.0 - 0.0)
X_train[0].min(), X_train[0].max()
```

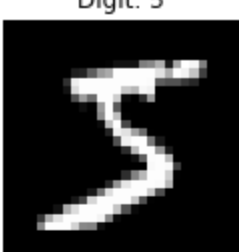
Out[5]:

```
(0.0, 1.0)
```

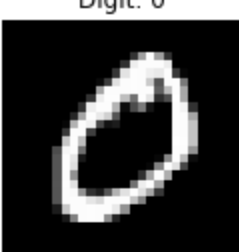
In [6]:

```
def plot_digit(image, digit, plt, i):
    plt.subplot(4, 5, i + 1)
    plt.imshow(image, cmap=plt.get_cmap('gray'))
    plt.title(f'Digit: {digit}')
    plt.xticks([])
    plt.yticks([])
plt.figure(figsize=(16, 10))
for i in range(20):
    plot_digit(X_train[i], y_train[i], plt, i)
plt.show()
```

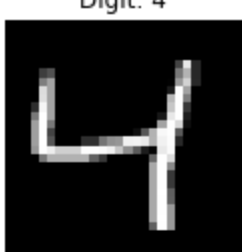
Digit: 5



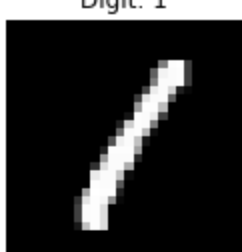
Digit: 0



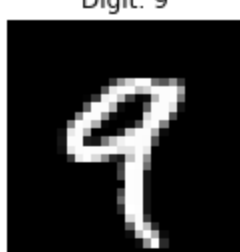
Digit: 4




Digit: 1



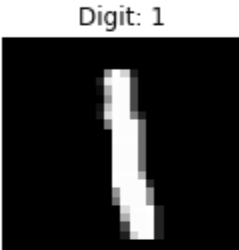
Digit: 9




Digit: 2



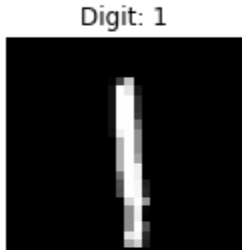
Digit: 1




Digit: 3



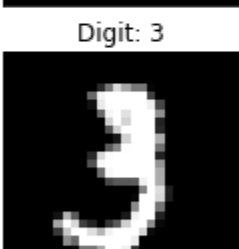
Digit: 1



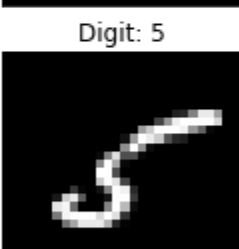
Digit: 4




Digit: 3




Digit: 5



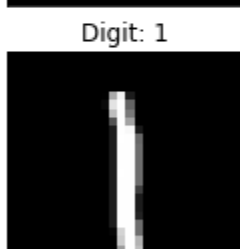
Digit: 3



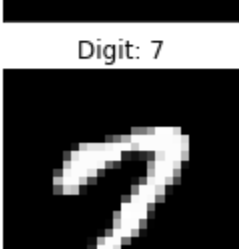
Digit: 6




Digit: 1



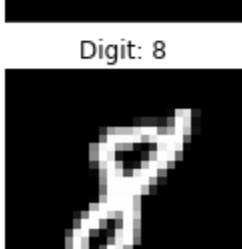
Digit: 7



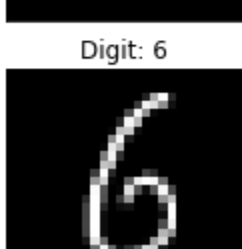
Digit: 2



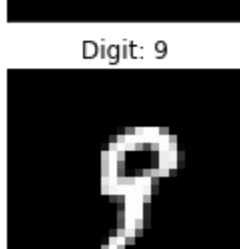
Digit: 8



Digit: 6



Digit: 9



In [7]:

```
X_train = X_train.reshape((X_train.shape + (1,)))
X_test = X_test.reshape((X_test.shape + (1,)))
```

In [8]:

```
y_train[0:20]
```

Out[8]:

```
array([5, 0, 4, 1, 9, 2, 1, 3, 1, 4, 3, 5, 3, 6, 1, 7, 2, 8, 6, 9],
      dtype=uint8)
```

In [9]:

```
model = Sequential([
    Conv2D(32, (3, 3), activation="relu", input_shape=(28, 28, 1)),
    MaxPooling2D((2, 2)),
    Flatten(),
    Dense(100, activation="relu"),
    Dense(10, activation="softmax")
])
```

In [10]:

```
optimizer = SGD(learning_rate=0.01, momentum=0.9)
model.compile(
    optimizer=optimizer,
    loss="sparse_categorical_crossentropy",
    metrics=["accuracy"]
)
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
flatten (Flatten)	(None, 5408)	0
dense (Dense)	(None, 100)	540900
dense_1 (Dense)	(None, 10)	1010

Total params: 542,230
Trainable params: 542,230
Non-trainable params: 0

In [11]:

```
model.fit(X_train, y_train, epochs=10, batch_size=32)
```

Out[11]:

```
<keras.callbacks.History at 0x27b921bdfa0>
```

In [12]:

```
plt.figure(figsize=(16, 10))
for i in range(20):
    image = random.choice(X_test).squeeze()
    digit = np.argmax(model.predict(image.reshape((1, 28, 28, 1))))[0], axis=-1)
    plot_digit(image, digit, plt, i)
plt.show()
```

1/1 [=====] - 0s 105ms/step

1/1 [=====] - 0s 19ms/step

1/1 [=====] - 0s 18ms/step

1/1 [=====] - 0s 19ms/step

1/1 [=====] - 0s 20ms/step

1/1 [=====] - 0s 20ms/step

1/1 [=====] - 0s 20ms/step

1/1 [=====] - 0s 18ms/step

1/1 [=====] - 0s 19ms/step

1/1 [=====] - 0s 19ms/step

1/1 [=====] - 0s 18ms/step

1/1 [=====] - 0s 19ms/step

1/1 [=====] - 0s 19ms/step

1/1 [=====] - 0s 18ms/step

1/1 [=====] - 0s 19ms/step

1/1 [=====] - 0s 19ms/step

1/1 [=====] - 0s 20ms/step

1/1 [=====] - 0s 20ms/step


1/1 [=====] - 0s 20ms/step

1/1 [=====] - 0s 20ms/step


1/1 [=====] - 0s 21ms/step

1/1 [=====] - 0s 20ms/step

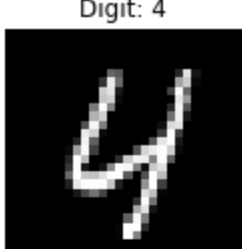
Digit: 8




Digit: 9




Digit: 4




Digit: 8




Digit: 7




Digit: 2




Digit: 0




Digit: 3



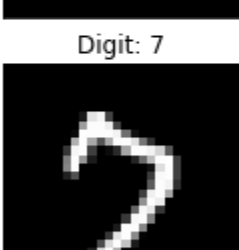
Digit: 8



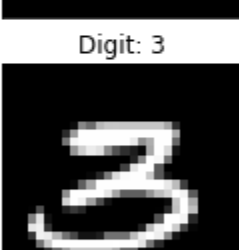
Digit: 8



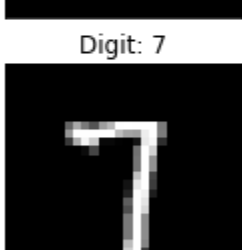
Digit: 7




Digit: 3




Digit: 7



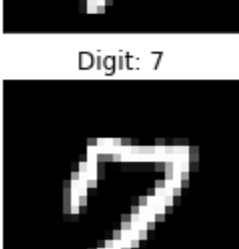
Digit: 9



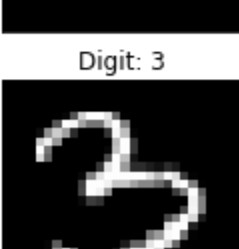
Digit: 6



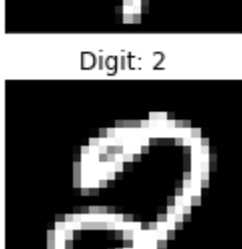
Digit: 7




Digit: 3



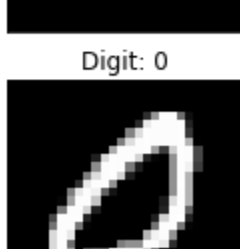
Digit: 2



Digit: 3



Digit: 0



In [13]:

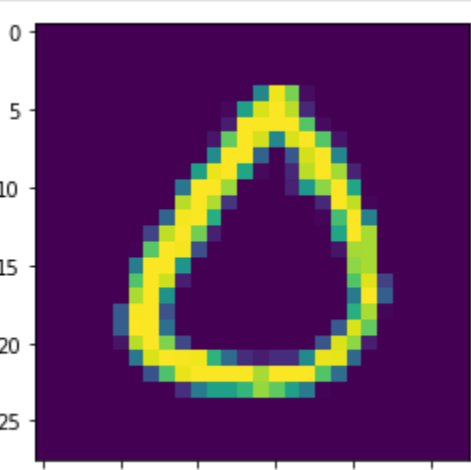
```
predictions = np.argmax(model.predict(X_test), axis=-1)
accuracy_score(y_test, predictions)
```

Out[13]:

```
313/313 [=====] - 1s 3ms/step
0.9881
```

In [14]:

```
n=random.randint(0,9999)
plt.imshow(X_test[n])
plt.show()
```



In [15]:

```
predicted_value=model.predict(X_test)
print("Handwritten number in the image is= %d"%np.argmax(predicted_value[n]))
```

Out[15]:

```
313/313 [=====] - 1s 3ms/step
Handwritten number in the image is= 0
```

In [16]:

```
score = model.evaluate(X_test, y_test, verbose=0)
print('Test loss:', score[0]) #Test loss: 0.0296396646054
print('Test accuracy:', score[1])
```

Out[16]:

```
Test loss: 0.03916610777378082
Test accuracy: 0.988099992275238
```

In [17]:

```
#The implemented CNN model is giving Loss=0.04624301567673683 and
#accuracy: 0.9872000217437744 for test mnist dataset
```


Name:Dipali Chavan

Roll No:14139

```
In [1]: import pandas as pd
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix, recall_score, accuracy_score, precision_score

RANDOM_SEED = 2021
TEST_PCT = 0.3
LABELS = ["Normal", "Fraud"]
```

```
In [2]: dataset = pd.read_csv("creditcard.csv")
```


```
In [3]: #check for any null values
print("Any nulls in the dataset",dataset.isnull().values.any())
print('-----')
print("No. of unique labels",len(dataset['Class'].unique()))
print("Label values",dataset.Class.unique())

#0 is for normal credit card transaction
#1 is for fraudulent credit card transaction
print('-----')
print("Break down of Normal and Fraud Transactions")
print(pd.value_counts(dataset['Class'],sort=True))
```

```
Any nulls in the dataset False
-----
No. of unique labels 2
Label values [0 1]
-----
Break down of Normal and Fraud Transactions
0    284315
1      492
Name: Class, dtype: int64
```


```
In [4]: #visualizing the imbalanced dataset
count_classes = pd.value_counts(dataset['Class'],sort=True)
count_classes.plot(kind='bar',rot=0)
plt.xticks(range(len(dataset['Class'].unique()))),dataset.Class.unique())
plt.title("Frequency by observation number")
plt.xlabel("Class")
plt.ylabel("Number of Observations")
```

Out[4]: Text(0, 0.5, 'Number of Observations')



```
In [4]: #Save the normal and fraudulent transactions in separate dataframe
normal_dataset = dataset[dataset.Class == 0]
fraud_dataset = dataset[dataset.Class == 1]

#Visualize transaction amounts for normal and fraudulent transactions
bins = np.linspace(200,2500,100)
plt.hist(normal_dataset.Amount,bins=bins,alpha=1,density=True,label='Normal')
plt.hist(fraud_dataset.Amount,bins=bins,alpha=0.5,density=True,label='Fraud')
plt.legend(loc='upper right')
plt.title("Transaction Amount vs Percentage of Transactions")
plt.xlabel("Transaction Amount (USD)")
plt.ylabel("Percentage of Transactions")
plt.show()
```



	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23	V24	V25	V26	V27	V28	Amount	Class
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.336321	0.462388	0.239599	0.096698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	0.128539	-0.189115	0.133558	-0.021053	149.62	0
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846	0.167170	0.125895	-0.008983	0.014724	2.69	0
2	1.0	-1.358354	-1.340163	1.73209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281	-0.327642	-0.139097	-0.055353	-0.059752	378.66	0
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575	0.647376	-0.221929	0.062723	0.061458	123.50	0
4	2.0	-1.582233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267	-0.206010	0.502292	0.219422	0.215153	69.99	0
...
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066556	-5.364473	-2.606837	-4.918215	7.305334	1.914428	...	0.213454	0.111864	1.014480	-0.509348	1.436807	0.250034	0.943651	0.823731	0.77	0
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.024330	0.294869	0.584800	...	0.214205	0.924384	0.012463	-1.016226	-0.606624	-0.395255	0.068472	-0.053527	24.79	0
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.296827	0.708417	0.432454	...	0.232045	0.578229	-0.037501	0.640134	0.265745	-0.087371	0.004455	-0.026561	67.88	0
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.686180	0.679145	0.392087	...	0.265245	0.800049	-0.163298	0.123205	-0.569159	0.546668	0.108821	0.104533	10.00	0
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.577006	-0.414650	0.486180	...	0.261057	0.643078	0.376777	0.008797	-0.473649	-0.818267	-0.002415	0.013649	217.00	0

284807 rows × 31 columns

```
In [6]: sc = StandardScaler()
dataset['Time'] = sc.fit_transform(dataset['Time'].values.reshape(-1,1))
dataset['Amount'] = sc.fit_transform(dataset['Amount'].values.reshape(-1,1))
```

```
In [7]: raw_data = dataset.values
#The last element contains if the transaction is normal which is represented by 0 and if fraud then 1
labels = raw_data[:, -1]

#The other data points are the electrocardiogram data
data = raw_data[:,0:-1]

train_data,test_data,train_labels,test_labels = train_test_split(data,labels,test_size = 0.2,random_state =2021)
```

```
In [8]: min_val = tf.reduce_min(train_data)
max_val = tf.reduce_max(train_data)

train_data = (train_data - min_val) / (max_val - min_val)
test_data = (test_data - min_val) / (max_val - min_val)

train_data = tf.cast(train_data,tf.float32)
test_data = tf.cast(test_data,tf.float32)
```

```
In [9]: train_labels = train_labels.astype(bool)
test_labels = test_labels.astype(bool)

#Creating normal and fraud datasets
normal_train_data = train_data[train_labels]
normal_test_data = test_data[test_labels]

fraud_train_data = train_data[~train_labels]
fraud_test_data = test_data[~test_labels]
print("No. of records in Fraud Train Data=",len(fraud_train_data))
print("No. of records in Normal Train Data=",len(normal_train_data))
print("No. of records in Fraud Test Data=",len(fraud_test_data))
print("No. of records in Normal Test Data=",len(normal_test_data))
```

```
No. of records in Fraud Train Data= 389
No. of records in Normal Train Data= 227456
No. of records in Fraud Test Data= 493
No. of records in Normal Test Data= 56859
```

```
In [10]: nb_epoch = 50
batch_size = 64
input_dim = normal_train_data.shape[1]
#num of columns,30
encoding_dim = 14
hidden_dim1 = int(encoding_dim / 2)
hidden_dim2 = 4
learning_rate = 1e-7
```

```
In [11]: #input layer
input_layer = tf.keras.layers.Input(shape=(input_dim,))

#Encoder
encoder = tf.keras.layers.Dense(encoding_dim,activation="tanh",activity_regularizer = tf.keras.regularizers.l2(learning_rate))(input_layer)
encoder = tf.keras.layers.Dropout(0.2)(encoder)
encoder = tf.keras.layers.Dense(hidden_dim1,activation='relu')(encoder)
encoder = tf.keras.layers.Dense(hidden_dim2,activation=tf.nn.leaky_relu)(encoder)

#Decoder
decoder = tf.keras.layers.Dense(hidden_dim1,activation='relu')(encoder)
decoder = tf.keras.layers.Dropout(0.2)(decoder)
decoder = tf.keras.layers.Dense(encoding_dim,activation='relu')(decoder)
decoder = tf.keras.layers.Dense(input_dim,activation='tanh')(decoder)

#Autoencoder
autoencoder = tf.keras.Model(inputs = input_layer,outputs = decoder)
autoencoder.summary()
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 30)]	0
dense (Dense)	(None, 14)	434
dropout (Dropout)	(None, 14)	0
dense_1 (Dense)	(None, 7)	105
dense_2 (Dense)	(None, 4)	32
dense_3 (Dense)	(None, 7)	35
dropout_1 (Dropout)	(None, 7)	0
dense_4 (Dense)	(None, 14)	112
dense_5 (Dense)	(None, 30)	450

=====

Total params: 1,168
Trainable params: 1,168
Non-trainable params: 0

```
In [12]: cp = tf.keras.callbacks.ModelCheckpoint(filepath="autoencoder_fraud.h5",mode="min",monitor='val_loss',verbose=2,save_best_only=True)
#Define our early stopping
early_stop = tf.keras.callbacks.EarlyStopping(
    monitor='val_loss',
    min_delta=0.0001,
    patience=10,
    verbose=1,
    mode='min',
    restore_best_weights=True
)
```


```
In [13]: autoencoder.compile(metrics=['accuracy'],loss= 'mean_squared_error',optimizer='adam')
```

```
In [14]: history = autoencoder.fit(normal_train_data,normal_train_data,epochs = nb_epoch,
                                batch_size = batch_size,shuffle = True,
                                validation_data = (test_data,test_data),
                                verbose=1,
                                callbacks = [cp,early_stop]).history
```

```
Epoch 1/50
3543/3554 [=====] - ETA: 0s - loss: 0.0033 - accuracy: 0.0372
Epoch 1: val_loss improved from inf to 0.00002, saving model to autoencoder_fraud.h5
3554/3554 [=====] - 11s 2ms/step - loss: 0.0033 - accuracy: 0.0372 - val_loss: 2.0179e-05 - val_accuracy: 0.0343
Epoch 2/50
3545/3554 [=====] - ETA: 0s - loss: 1.9583e-05 - accuracy: 0.0609
Epoch 2: val_loss did not improve from 0.00002
3554/3554 [=====] - 6s 2ms/step - loss: 1.9599e-05 - accuracy: 0.0609 - val_loss: 2.0190e-05 - val_accuracy: 0.0078
Epoch 3/50
3517/3554 [=====] - ETA: 0s - loss: 1.9562e-05 - accuracy: 0.0619
Epoch 3: val_loss improved from 0.00002 to 0.00002, saving model to autoencoder_fraud.h5
3554/3554 [=====] - 7s 2ms/step - loss: 1.9580e-05 - accuracy: 0.0619 - val_loss: 2.0025e-05 - val_accuracy: 0.0420
Epoch 4/50
3527/3554 [=====] - ETA: 0s - loss: 1.9545e-05 - accuracy: 0.0599
Epoch 4: val_loss did not improve from 0.00002
3554/3554 [=====] - 6s 2ms/step - loss: 1.9530e-05 - accuracy: 0.0601 - val_loss: 2.0277e-05 - val_accuracy: 0.2168
Epoch 5/50
3549/3554 [=====] - ETA: 0s - loss: 1.8826e-05 - accuracy: 0.1758
Epoch 5: val_loss improved from 0.00002 to 0.00002, saving model to autoencoder_fraud.h5
3554/3554 [=====] - 6s 2ms/step - loss: 1.8831e-05 - accuracy: 0.1759 - val_loss: 1.8344e-05 - val_accuracy: 0.2184
Epoch 6/50
3539/3554 [=====] - ETA: 0s - loss: 1.7526e-05 - accuracy: 0.2362
Epoch 6: val_loss improved from 0.00002 to 0.00002, saving model to autoencoder_fraud.h5
3554/3554 [=====] - 6s 2ms/step - loss: 1.7518e-05 - accuracy: 0.2363 - val_loss: 1.7095e-05 - val_accuracy: 0.3538
Epoch 7/50
3516/3554 [=====] - ETA: 0s - loss: 1.8826e-05 - accuracy: 0.1125
Epoch 7: val_loss did not improve from 0.00002
3554/3554 [=====] - 5s 2ms/step - loss: 1.8813e-05 - accuracy: 0.1137 - val_loss: 1.7990e-05 - val_accuracy: 0.2041
Epoch 8/50
3529/3554 [=====] - ETA: 0s - loss: 1.7206e-05 - accuracy: 0.2141
Epoch 8: val_loss did not improve from 0.00002
3554/3554 [=====] - 5s 1ms/step - loss: 1.7328e-05 - accuracy: 0.2143 - val_loss: 1.7153e-05 - val_accuracy: 0.2781
Epoch 9/50
3518/3554 [=====] - ETA: 0s - loss: 1.6837e-05 - accuracy: 0.2481
Epoch 9: val_loss improved from 0.00002 to 0.00002, saving model to autoencoder_fraud.h5
3554/3554 [=====] - 6s 2ms/step - loss: 1.6825e-05 - accuracy: 0.2482 - val_loss: 1.6811e-05 - val_accuracy: 0.3518
Epoch 10/50
3516/3554 [=====] - ETA: 0s - loss: 1.6692e-05 - accuracy: 0.2509
Epoch 10: val_loss improved from 0.00002 to 0.00002, saving model to autoencoder_fraud.h5
3554/3554 [=====] - 6s 2ms/step - loss: 1.6682e-05 - accuracy: 0.2510 - val_loss: 1.6469e-05 - val_accuracy: 0.3492
Epoch 11/50
3526/3554 [=====] - ETA: 0s - loss: 1.6569e-05 - accuracy: 0.2484
Epoch 11: val_loss improved from 0.00002 to 0.00002, saving model to autoencoder_fraud.h5
Restoring model weights from the end of the best epoch: 1.
3554/3554 [=====] - 7s 2ms/step - loss: 1.6561e-05 - accuracy: 0.2484 - val_loss: 1.6237e-05 - val_accuracy: 0.2865
Epoch 11: early stopping
```

```
In [15]: plt.plot(history['loss'],linewidth = 2,label = 'Train')
plt.plot(history['val_loss'],linewidth = 2,label = 'Test')
plt.legend(loc='upper right')
plt.title('Model Loss')
plt.xlabel('Loss')
plt.ylabel('Epoch')

#plt.ylim(ymin=0.70,ymax=1)
```




```
In [16]: test_x_predictions = autoencoder.predict(test_data)
mse = np.mean(np.power(test_data - test_x_predictions, 2),axis = 1)
error_df = pd.DataFrame({'Reconstruction_error':mse,
                        'True_class':test_labels})
```

1781/1781 [=====] - 2s 865us/step

```
In [17]: threshold_fixed = 50
groups = error_df.groupby('True_class')
fig,ax = plt.subplots()


for name,group in groups:
    ax.plot(group.index,group.Reconstruction_error,markers='o',ms = 3.5,linestyle='',
            label = "Fraud" if name==1 else "Normal")
ax.lines(threshold_fixed,ax.get_xlim()[0],ax.get_xlim()[2],colors="r",zorder=100,label="Threshold")
ax.legend()
plt.title("Reconstructions error for normal and fraud data")
plt.xlabel("Reconstruction error")
plt.ylabel("Data point index")
plt.show()
```



```
In [19]: threshold_fixed = 52
pred_y = [1 if e > threshold_fixed else 0
          for e in
          error_df.Reconstruction_error.values]
error_df['pred'] = pred_y
conf_matrix = confusion_matrix(error_df.True_class,pred_y)
```

```
plt.figure(figsize = (4,4))
sns.heatmap(conf_matrix,xticklabels = LABELS,yticklabels = LABELS,annot = True,fmt="d")
plt.title("Confusion matrix")
plt.xlabel("True class")
plt.ylabel("Predicted class")
plt.show()

#Print Accuracy,Precision and Recall
print("Accuracy :",accuracy_score(error_df['True_class'],error_df['pred']))
print("Recall :",recall_score(error_df['True_class'],error_df['pred']))
print("Precision :",precision_score(error_df['True_class'],error_df['pred']))
```



Accuracy : 0.9981917769741231
Recall : 0.0
Precision : 0.0

C:\Users\Manish\conda\envs\tensorflow\lib\site-packages\sklearn\metrics_classification.py:1318: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to no pr
dicted samples. Use 'zero_division' parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))

In []:

Name:Dipali Chavan
Roll No : I4139

In [1]:

```
import matplotlib.pyplot as plt
import seaborn as sns
import matplotlib as mpl
import matplotlib.pyplot as pylab
import numpy as np
%matplotlib inline
```

In [2]:

```
#Data Prepration
import re
```

In [3]:

```
sentences = """We are about to study the idea of a computational process.
Computational processes are abstract beings that inhabit computers.
As they evolve, processes manipulate other abstract things called data.
The evolution of a process is directed by a pattern of rules
called a program. People create programs to direct processes. In effect,
we conjure the spirits of the computer with our spells."""
```

Clean Data

In [4]:

```
# remove special characters
sentences = re.sub('[^A-Za-z0-9]+', ' ', sentences)

# remove 1 letter words
sentences = re.sub(r'(?<^>)\w(?:$| )', ' ', sentences).strip()

# lower all characters
sentences = sentences.lower()
```

Vocabulary

In [5]:

```
words = sentences.split()
vocab = set(words)
```

In [6]:

```
vocab_size = len(vocab)
embed_dim = 10
context_size = 2
```

Implementation

In [7]:

```
word_to_ix = {word: i for i, word in enumerate(vocab)}
ix_to_word = {i: word for i, word in enumerate(vocab)}
```

Data bags

In [8]:

```
# data - [(context), target]

data = []
for i in range(2, len(words) - 2):
    [words[i - 2], words[i - 1], words[i + 1], words[i + 2]]
```

```
target = words[i]
data.append((context, target))
print(data[:5])
```

```
[[('we', 'are', 'to', 'study'], 'about'), ('are', 'about', 'study', 'the'], 'to'), ('about', 'to', 'the', 'idea'], 'study'), ('to', 'study', 'idea', 'of'], 'the'), ('study', 'the', 'of', 'computational'], 'idea')]
```

Embeddings

```
In [9]: embeddings = np.random.random_sample((vocab_size, embed_dim))
```

Linear Model

```
In [10]: def linear(m, theta):
          w = theta
          return m.dot(w)
```

Log softmax + NLLloss = Cross Entropy

```
In [11]: def log_softmax(x):
          e_x = np.exp(x - np.max(x))
          return np.log(e_x / e_x.sum())
```

```
In [12]: def NLLLoss(logs, targets):
          out = logs[range(len(targets)), targets]
          return -out.sum()/len(out)
```

```
In [13]: def log_softmax_crossentropy_with_logits(logits, target):

          out = np.zeros_like(logits)
          out[np.arange(len(logits)), target] = 1

          softmax = np.exp(logits) / np.exp(logits).sum(axis=-1, keepdims=True)

          return (- out + softmax) / logits.shape[0]
```

Forward function

```
In [14]: def forward(context_idxs, theta):
          m = embeddings[context_idxs].reshape(1, -1)
          n = linear(m, theta)
          o = log_softmax(n)

          return m, n, o
```

Backward function

```
In [15]: def backward(preds, theta, target_idxs):
          m, n, o = preds

          dlog = log_softmax_crossentropy_with_logits(n, target_idxs)
          dw = m.T.dot(dlog)

          return dw
```

Optimize function

```
In [16]: def optimize(theta, grad, lr=0.03):  
        theta -= grad * lr  
        return theta
```

Training

```
In [17]: #Generate training data  
  
theta = np.random.uniform(-1, 1, (2 * context_size * embed_dim, vocab_size))
```

```
In [18]: epoch_losses = {}  
  
for epoch in range(80):  
    losses = []  
  
    for context, target in data:  
        context_idx = np.array([word_to_ix[w] for w in context])  
        preds = forward(context_idx, theta)  
  
        target_idx = np.array([word_to_ix[target]])  
        loss = NLLLoss(preds[-1], target_idx)  
  
        losses.append(loss)  
  
        grad = backward(preds, theta, target_idx)  
        theta = optimize(theta, grad, lr=0.03)  
  
    epoch_losses[epoch] = losses
```


Analyze

Plot loss/epoch

```
In [19]: ix = np.arange(0, 80)  
  
fig = plt.figure()  
fig.suptitle('Epoch/Losses', fontsize=20)  
plt.plot(ix, [epoch_losses[i][0] for i in ix])  
plt.xlabel('Epochs', fontsize=12)  
plt.ylabel('Losses', fontsize=12)
```

```
Out[19]: Text(0, 0.5, 'Losses')
```


Epoch/Losses



Predict function

```
In [20]: def predict(words):
         context_idxes = np.array([word_to_ix[w] for w in words])
         preds = forward(context_idxes, theta)
         word = ix_to_word[np.argmax(preds[-1])]

         return word
```

```
In [21]: # (['we', 'are', 'to', 'study'], 'about')
         predict(['we', 'are', 'to', 'study'])
```

Out[21]: 'about'

Accuracy

```
In [22]: def accuracy():
         wrong = 0

         for context, target in data:
             if(predict(context) != target):
                 wrong += 1

         return (1 - (wrong / len(data)))
```

```
In [23]: accuracy()
```

Out[23]: 1.0

```
In [24]: predict(['processes', 'manipulate', 'things', 'study'])
```

Out[24]: 'effect'

```
In [ ]:
```

In []:

Name : Dipali Chavan
Roll No : I4139

In [1]:

```
from tensorflow.keras.preprocessing.image import load_img
from tensorflow.keras.preprocessing.image import img_to_array
from keras.applications.vgg16 import preprocess_input
from keras.applications.vgg16 import decode_predictions
from keras.applications.vgg16 import VGG16
# load an image from file
image = load_img('download.jpg', target_size=(224, 224))
# convert the image pixels to a numpy array
image = img_to_array(image)
# reshape data for the model
image = image.reshape((1, image.shape[0], image.shape[1], image.shape[2]))
# prepare the image for the VGG model
image = preprocess_input(image)
# load the model
model = VGG16()
# predict the probability across all output classes
yhat = model.predict(image)
# convert the probabilities to class labels
label = decode_predictions(yhat)
# retrieve the most likely result, e.g. highest probability
label = label[0][0]
# print the classification
print('%s (%.2f%%)' % (label[1], label[2]*100))
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels.h5
553467096/553467096 [=====] - 151s 0us/step
1/1 [=====] - 2s 2s/step
Downloading data from https://storage.googleapis.com/download.tensorflow.org/data/imagenet_class_index.json
35363/35363 [=====] - 0s 0us/step
castle (34.03%)

In [2]:

```
# load an image from file
image = load_img('download2.png', target_size=(224, 224))
# convert the image pixels to a numpy array
image = img_to_array(image)
# reshape data for the model
image = image.reshape((1, image.shape[0], image.shape[1], image.shape[2]))
# prepare the image for the VGG model
image = preprocess_input(image)
# load the model
model = VGG16()
# predict the probability across all output classes
yhat = model.predict(image)
# convert the probabilities to class labels
label = decode_predictions(yhat)
# retrieve the most likely result, e.g. highest probability
label = label[0][0]
# print the classification
print('%s (%.2f%%)' % (label[1], label[2]*100))
```

1/1 [=====] - 1s 672ms/step
valley (44.85%)

In [3]:

```
# load an image from file
img('download3.jpg', target_size=(224, 224))
```



```
# convert the image pixels to a numpy array
image = img_to_array(image)
# reshape data for the model
image = image.reshape((1, image.shape[0], image.shape[1], image.shape[2]))
# prepare the image for the VGG model
image = preprocess_input(image)
# load the model
model = VGG16()
# predict the probability across all output classes
yhat = model.predict(image)
# convert the probabilities to class labels
label = decode_predictions(yhat)
# retrieve the most likely result, e.g. highest probability
label = label[0][0]
# print the classification
print('%s (%.2f%%)' % (label[1], label[2]*100))
```

```
1/1 [=====] - 2s 2s/step
golden_retriever (84.78%)
```

In []: